


What the model says vs what it exhibits



*Activation lens: first calibration numbers (09/26).
Consumer-metric gating catches 19-21 pp more real behavior
changes than raw L2 at matched flag rates. Shakedown result;
replication running. Probe mid-layers: the final-layer $\hat{p}$ map
fails held-out validation.*

turboquant-pro (PyPI): PyTorch-native compression
HF 1-liner: TurboQuantCache() · Triton + Volta kernels · vLLM ~5× KV
live here: 3-bit codec on NATS · claims CI-gated · measured false-clear diag (v2.0.0a3)