


What the model says vs what it exhibits



*activation / probe lens is EARLY —
uncalibrated by default (research-grade)*

turboquant-pro (PyPI): PyTorch-native compression
HF 1-liner: `generate(past_key_values=TurboQuantCache())` · Triton + Volta kernels · vLLM ~5× KV
live in this stack: 3-bit embedding codec on the NATS memory bus · claims CI-gated (CLAIMS.md)