


What the model says vs what it exhibits



*activation / probe lens is EARLY —
uncalibrated by default (research-grade)*

turboquant-pro (PyPI): PyTorch-native compression
HF 1-liner: `TurboQuantCache()` · Triton + Volta kernels · vLLM ~5× KV
live here: 3-bit embedding codec on NATS · claims CI-gated (CLAIMS.md)